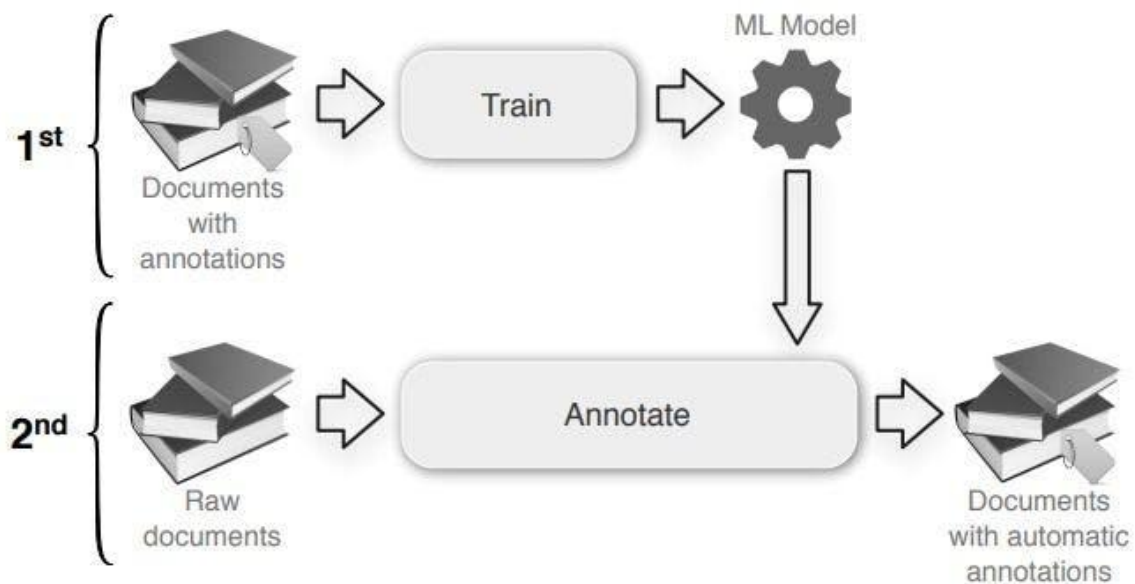


## Introduction

Named Entity Recognition (NER) is a crucial NLP technique that allows machines to **automatically identify and categorize important elements in text**, such as names, dates, organizations, and locations. From my perspective, NER acts like a smart highlighter—it transforms unstructured text into structured insights, making it easier to analyze and act upon. Using advanced models like **spaCy's transformer-based pipeline**, we can achieve high accuracy in recognizing entities while also providing flexibility to handle different text types and domains. This makes NER invaluable for tasks like information extraction, data analysis, and automated content understanding.



## Installing PIP & Importing The Files

```
0s # Install required packages if not already installed
# pip install spacy spacy-transformers pandas matplotlib
# python -m spacy download en_core_web_trf

import spacy
from spacy import displacy
import pandas as pd
from collections import Counter
import matplotlib.pyplot as plt

[3] pip install torch

Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.8.0+cu126)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.19.1)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch) (4.14.1)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch) (75.2.0)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch) (1.13.3)
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages (from torch) (3.5)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.12/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages (from torch) (2025.3.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.80)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch) (9.10.2.21)
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.4.1)
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch) (11.3.0.4)
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch) (10.3.7.77)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch) (11.7.1.2)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch) (12.5.4.2)
Requirement already satisfied: nvidia-cusparselt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch) (0.7.1)
Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in /usr/local/lib/python3.12/dist-packages (from torch) (2.27.3)
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.85)
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch) (1.11.1.6)
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-packages (from torch) (3.4.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from Jinja2->torch) (3.0.2)
```

## Loading NLP Model /Input Text/Process Text

```

# Step 1: Load NLP Model Safely
def load_model():
    try:
        # Try loading transformer-based model first
        nlp = spacy.load("en_core_web_trf")
        print("Loaded transformer-based model: en_core_web_trf")
    except (OSError, ValueError) as e:
        print("Transformer model not available. Falling back to small model.")
        print("Error:", e)
    try:
        nlp = spacy.load("en_core_web_sm")
        print("Loaded small English model: en_core_web_sm")
    except OSError:
        print("Small model not found. Downloading...")
        import subprocess
        subprocess.run(["python", "-m", "spacy", "download", "en_core_web_sm"])
        nlp = spacy.load("en_core_web_sm")
    return nlp
nlp = load_model()

```

Transformer model not available. Falling back to small model.  
 Error: [E002] Can't find factory for 'curated\_transformer' for language English (en). This usually happens when you have a custom tokenizer or a custom parser.  
 Available factories: merge\_noun\_chunks, merge\_entities, merge\_subtokens, en.lemmatizer  
 Loaded small English model: en\_core\_web\_sm

```

[12] # Step 2: Input Text
# -----
text = """
Apple Inc. announced a new iPhone on September 12, 2025 in Cupertino, California.
Elon Musk visited New York City on 15th August 2025 to discuss SpaceX partnerships.
Microsoft Corp. and Google LLC signed a contract on January 10, 2025.
"""

```

```

[13] # Step 3: Process Text

doc = nlp(text)

```

## Extract Entities & Count Entity Types

```
# Step 4: Extract Entities
entities = [(ent.text, ent.label_) for ent in doc.ents]
if not entities:
    print("No entities found.")
else:
    df_entities = pd.DataFrame(entities, columns=["Entity", "Type"])
    print("Extracted Entities:")
    print(df_entities)
```

	Entity	Type
0	Apple Inc.	ORG
1	September 12, 2025	DATE
2	Cupertino	GPE
3	California	GPE
4	Elon Musk	PERSON
5	New York City	GPE
6	15th	ORDINAL
7	August 2025	DATE
8	Microsoft Corp.	ORG
9	Google LLC	ORG
10	January 10, 2025	DATE

```
[15] # Step 5: Count Entity Types
entity_counter = Counter([ent.label_ for ent in doc.ents])
print("\nEntity Counts:")
print(entity_counter)
```


```
Entity Counts:
Counter({'ORG': 3, 'DATE': 3, 'GPE': 3, 'PERSON': 1, 'ORDINAL': 1})
```

## Visualize the entities

```
# Step 6: Visualize Entities

# Inline visualization (Jupyter Notebook)
try:
    displacy.render(doc, style="ent", jupyter=True)
except:
    print("Displacy render works only in Jupyter Notebook.")
```

```
[17] # Bar chart
plt.figure(figsize=(8, 4))
plt.bar(entity_counter.keys(), entity_counter.values(), color="skyblue")
plt.title("Entity Type Counts")
plt.xlabel("Entity Type")
plt.ylabel("Count")
plt.show()
```



```
# Step 7: Filter by Specific Types
people = [ent.text for ent in doc.ents if ent.label_ == "PERSON"]
dates = [ent.text for ent in doc.ents if ent.label_ == "DATE"]
companies = [ent.text for ent in doc.ents if ent.label_ in ["ORG", "COMPANY"]]

print("\nPeople:", people)
print("Dates:", dates)
print("Companies:", companies)
```



```
People: ['Elon Musk']
Dates: ['September 12, 2025', 'August 2025', 'January 10, 2025']
Companies: ['Apple Inc.', 'Microsoft Corp.', 'Google LLC']
```

```
[19] # Step 8: Export to CSV
try:
    df_entities.to_csv("ner_output.csv", index=False)
    print("Entities exported to ner_output.csv")
except Exception as e:
    print("Error saving CSV:", e)
```



```
Entities exported to ner_output.csv
```

## **Conclusion**

From my perspective, NER is not just about labeling words—it's about unlocking the hidden structure within text. By leveraging a robust NLP pipeline, we can quickly extract actionable information, visualize patterns, and organize data efficiently. While transformer models offer superior accuracy, it's important to implement fallback mechanisms for reliability across different environments. Overall, NER provides a bridge between raw textual data and meaningful insights, empowering applications that range from business intelligence to AI-driven analytics.